

The background of the slide is a dark blue and teal gradient. It features several semi-transparent financial charts, including a candlestick chart and a bar chart, overlaid with binary code (0s and 1s) and various numerical data points. The overall aesthetic is high-tech and data-driven.

# SuperMUC-NG

---

Francesc Contreras & Albert Pérez



# Content

---

- Why SuperMUC-NG ?
- Specifications & Analysis
- Comparison
- Conclusions



# Why SuperMUC-NG ?

---

- Based on Lenovo and manufactured in Germany
- Energy efficiency
- Singular refrigeration: Hot 🔥 Water to cool the nodes
- Xeon Platinum & Intel Omni-Path
- “Meta-Heuristics for Scheduling in Cluster Federated Env.”

Thesis by Eloi Gabaldon



# Specifications & Analysis

Manufactured by Lenovo and located in Garching, Germany.

Processor: Xeon Platinum 8174 24C  
3.1GHz

Internal Interconnection: Omni-Path  
network of 100 Gbit/s

Cooler System: Direct Warm Water Cooling

Total Compute Nodes: 9 Islands, 6.480  
Nodes, 311,040 Cores

- Thin Nodes: 8 Islands, 6.336 Nodes, 304.128 Cores
- Fat Nodes: 1 Island, 144 Nodes, 6.912 Cores

Memory (TByte): 719

# Nodes

---

## 144 Fat nodes

- Single node performance
- Ultra-high reliability
- High memory capacity

## 6336 Thin nodes

- Inexpensive
- High-speed CPU
- Well-suited for sequential workloads

# XEON Platinum

---

## XEON Platinum 8174

- 24C 3.10Ghz
- 33MB Cache DDR4-2666
- 6 Memory Channels
- Scalability s8s 3 QPI-Links
- UPI Links
- PCI Express 3.0



# Intel Omni-Path

---

## Intel Omni-Path


- High-performance communication
  - Low latency
  - Low power consumption
  - High throughput
- For small and large scale



# Intel OPA Performance

---

Intel OPA to Enhanced Data Rate (EDR) InfiniBand (IF)  
Ohio State University (OSU) OMB osu\_latency benchmark



- 8 bytes message
- 11% lower




# Intel OPA Performance

---


Ohio State University (OSU) OMB osu\_bw benchmark

- 1MB messages
- Both full wire rate of 100Gbps



# Intel OPA Performance

Ohio State University (OSU) OMB osu\_mbw\_mr benchmark



- 32 MPI rank pairs
- 64% higher


# Comparison

---

- With the first 4 current supercomputers
  - Fugaku, Summit, Sierra, and Sunway TaihuLight
- Based on
  - Cores
  - $R_{peak}$  /  $R_{max}$
  - $N_{max}$

# Comparison - Cores


---



- Clear difference
- But, it really affects directly the performance?

# Comparison - Rpeak

---




- LINPACK benchmark
- Theoretical peak performance



# Comparison - Rmax


---



- LINPACK benchmark
- Maximal achieved performance

# Comparison - Nmax

---




- Problem size
- To achieve Rmax

# Cooling System

---

- Warm water
- 97% heat removal efficiency
- Aquasar by IBM
- Water 4000 times efficient than air
- The waste heat produced is recycled
- Renewable energy sources



<https://www.youtube.com/watch?v=FbGyAXsLzIc>

# CONCLUSIONS

---

- Deepen in the Supercomputers' world
- Supercomputer vs Home Computer
- Architectures
- Energy Efficiency
- Cooling systems

Many thanks

---

Any questions?